

DLHLP - HW1


End-to-end Speech Recognition

TA: 莊永松、柯上優
dlhlp.ta@gmail.com

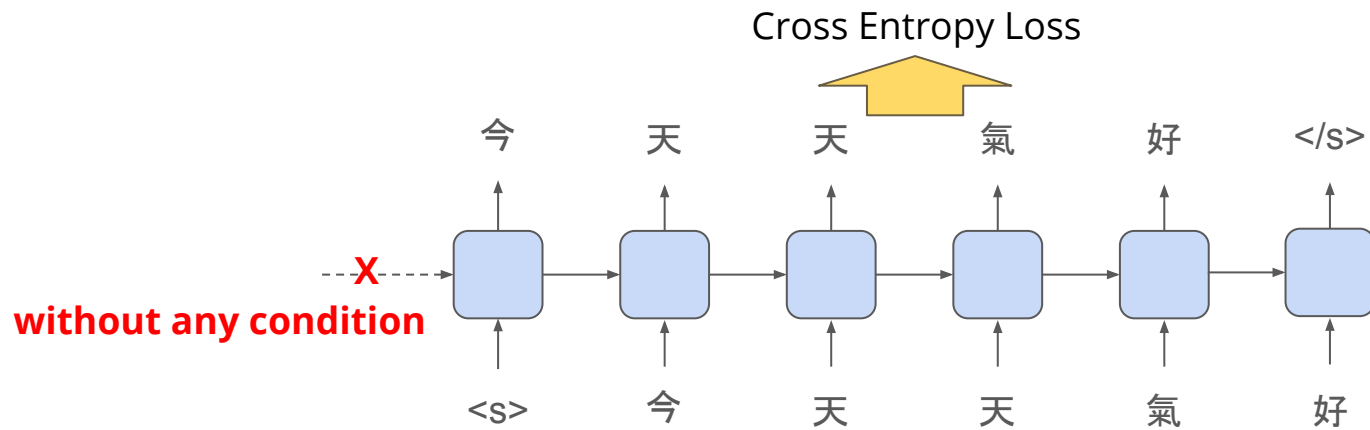
Seq2Seq attention-based ASR




Joint CTC-attention based ASR



RNN-LM (for re-scoring)



Dataset



To make HW1 easier...

- We use a small 10-hours Mandarin Chinese audio corpus
- All from a single speaker (no speaker variant problem)
- We transform text to ㄅ ㄆ ㄇ ㄏ to make vocabulary size smaller and help the model to converge faster
- You would not need to train a “real-world” ASR model
- About **2~4 hours** to converge on K80, while real-world ASR may need weeks to converge

Download Data

- Download Data(2.88G)

- <https://drive.google.com/file/d/1daFU8tPPUyhN7Fc6JUTohEfHXIn6ZDgq/view?usp=sharing>
- train: 8000, dev: 1000, test: 1000
- the transcript file in testing set is not removed for convenience, but all the answer is replaced with ㄅ ㄆ ㄇ ㄏ...

Hint: In Linux, use `bash get_dataset.sh`

https://github.com/DLHLP2020/hw1-speech-recognition/blob/master/get_dataset.sh

```
/DLHLP
├── /train
│   ├── 000001.wav
│   ├── .....
│   ├── 008000.wav
│   └── bopomo.trans.txt
├── /dev
│   ├── 008001.wav
│   ├── .....
│   ├── 009000.wav
│   └── bopomo.trans.txt
├── /test
│   ├── 009001.wav
│   ├── .....
│   ├── 010000.wav
│   └── bopomo.trans.txt
└── text-data.txt
```


Extract features from audio file w/ torchaudio

Step 1. Load WAV File

```
waveform, sample_freq = torchaudio.load("000001.wav")
```

waveform.shape: (1, 116400)

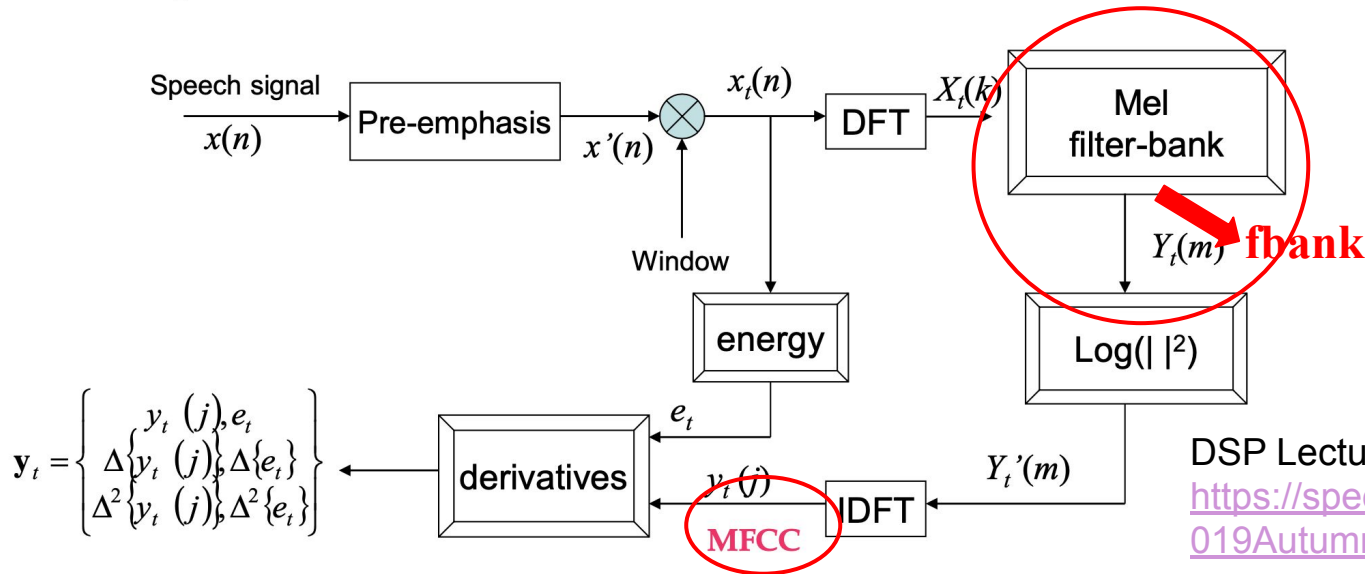
sample_freq: 48000



Step 2. Get fbank or MFCC from waveform

- **Mel-Frequency Cepstral Coefficients (MFCC)**

- Most widely used in the speech recognition
- Has generally obtained a better accuracy at relatively low computational complexity
- The process of MFCC extraction :



DSP Lecture 7:

<https://speech.ee.ntu.edu.tw/DSP2019Autumn/Slides/7.0.pptx>

Step 2. Get Fbank or MFCC from waveform

MFCC

```
feature = torchaudio.compliance.kaldi.mfcc(waveform,  
                                             sample_frequency=sample_freq)
```


feature.shape: (241, 13)

Fbank

```
feature = torchaudio.compliance.kaldi.fbank(waveform,  
                                             sample_frequency=sample_freq,  
                                             num_mel_bins=40)
```

feature.shape: (241, 40)

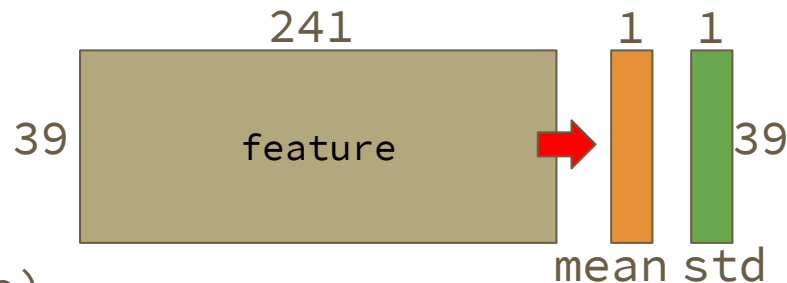
Step 3. Add Deltas



```
d1 = torchaudio.functional.compute_deltas(feature)
d2 = torchaudio.functional.compute_deltas(d1)
feature = torch.cat([feature, d1, d2], dim=-1)
feature.shape: (241, 39) or (241, 120)
```

Step 4. CMVN (Normalization)

```
eps = 1e-10
mean = feature.mean(0, keepdim=True)
std = feature.std(0, keepdim=True)
feature = (feature - mean) / (std + eps)
```



Sample code (strongly recommended)

<https://github.com/Alexander-H-Liu/End-to-end-ASR-Pytorch>

A completed & stable ASR implementation (Thanks to 劉浩然學長!)

Please read README before using it!

Alexander-H-Liu / End-to-end-ASR-Pytorch

Watch 29 Unstar 472 Fork 143

Code Issues 11 Pull requests 1 Actions Projects 0 Wiki Security Insights

This is an open source project (formerly named Listen, Attend and Spell - PyTorch Implementation) for end-to-end ASR implemented with Pytorch, the well known deep learning toolkit.

73 commits 3 branches 0 packages 0 releases 3 contributors MIT

Branch: master New pull request

Create new file Upload files Find file Clone or download

Alexander-H-Liu	fix adadelta init for testing & char decoder repeat issue	Latest commit fa683b 28 days ago
bin	fix adadelta init for testing & char decoder repeat issue	28 days ago
config	hybrid typo fixed	2 months ago
corpus	Fixed scheduler, corpus loading assertion, default lm	4 months ago
src	fix adadelta init for testing & char decoder repeat issue	28 days ago
tests	rename subword 16k	4 months ago
util	readme draft, libri example	4 months ago
.gitignore	add *.swp	7 months ago
LICENSE	Rebase to clear history	12 months ago



End-to-end Automatic Speech Recognition Systems - PyTorch Implementation

This is an open source project (formerly named Listen, Attend and Spell - PyTorch Implementation) for end-to-end ASR by Tzu-Wei Sung and me. Implementation was mostly done with Pytorch, the well known deep learning toolkit.

The end-to-end ASR was based on Listen, Attend and Spell¹. Multiple techniques proposed recently were also implemented, serving as additional plug-ins for better performance. For the list of techniques implemented, please refer to the [highlights](#), [configuration](#) and [references](#).

Feel free to use/modify them, any bug report or improvement suggestion will be appreciated. If you find this project helpful for your research, please do consider to cite [our paper](#), thanks!

Highlights



TODO: Add new Dataset

- See the example [corpus/librispeech.py](#)
- Create a copy [corpus/dlhlpy.py](#). Replace all “Libri” to “Dlhlpy” in the file.
 1. open each *.flac files → open each *.wav files
 2. def read_text()
 - origin: src_file = '-'.join(file.split('-')[:-1])+'.trans.txt'
 - ours: src_file = file.rsplit('/', 1)[0]+'bopomo.trans.txt'
 - Base on the format of our dataset

```
28 class DlhlpyDataset(Dataset):
29     def __init__(self, path, split, tokenizer, bucket_size, ascending=False):
30         # Setup
31         self.path = path
32         self.bucket_size = bucket_size
33
34         # List all wave files
35         file_list = []
36         for s in split:
37             split_list = list(Path(join(path, s)).rglob("*.wav"))
38             assert len(split_list) > 0, "No data found @ {}".format(join(path,s))
39             file_list += split_list
```

```
15 def read_text(file):
16     """Get transcription of target wave file,
17     it's somewhat redundant for accessing each txt multiplt times,
18     but it works fine with multi-thread"""
19     src_file = file.rsplit('/', 1)[0]+'bopomo.trans.txt'
20     idx = file.split('/')[-1].split('.')[0]
21
22     with open(src_file, 'r') as fp:
23         for line in fp:
24             if idx == line.split(' ')[0]:
25                 return line[:-1].split(' ', 1)[1]
26
```

TODO: Import new Dataset

- In the [src/data.py](#)
 - def create_dataset
 - def create_textset

```
64 def create_dataset(tokenizer, ascending, name, path, bucketing, batch_size,
65                    train_split=None, dev_split=None, test_split=None):
66     """ Interface for creating all kinds of dataset"""
67
68     # Recognize corpus
69     if name.lower() == "librispeech":
70         from corpus.librispeech import LibriDataset as Dataset
71     elif name.lower() == "dlhlp":
72         from corpus.dlhlp import DlhlpDataset as Dataset
73     else:
74         raise NotImplementedError
75
```

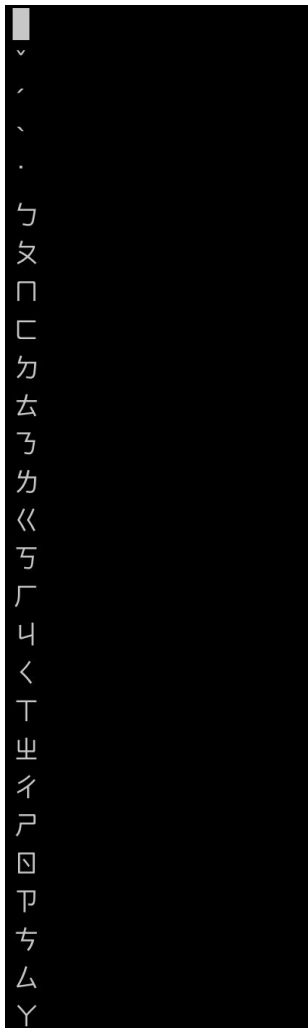
```
111 def create_textset(tokenizer, train_split, dev_split, name, path, buck
112                    """ Interface for creating all kinds of text dataset"""
113     msg_list = []
114
115     # Recognize corpus
116     if name.lower() == "librispeech":
117         from corpus.librispeech import LibriTextDataset as Dataset
118     elif name.lower() == "dlhlp":
119         from corpus.dlhlp import DlhlpTextDataset as Dataset
120     else:
121         raise NotImplementedError

```

TODO: Prepare vocab-count file

```
$ python util/generate_vocab_file.py  
    --input_file <input_text_file>  
    --mode character  
    --output_file <output_vocab_file>
```

we provide “text-data.txt” in the dataset as input_text file.



TODO: Write your own config file (1/3)

- See example [config/libri/asr_example.yaml](#)
- Create copys as you want to train a new model (e.g. asr_dlhlp.yaml)
 - path: where you unzip the data (e.g. 'data/DLHLP')
 - train/dev_split: dir name for train/dev under 'path'
 - audio feature setting: follow the original setting
 - mode: 'character'
 - vocab_file: the vocab file your have prepared

```
1 data:
2   corpus:
3     name: 'DLhlp'
4     path: 'data/DLHLP'
5     train_split: ['train']
6     dev_split: ['dev']
7     bucketing: True
8     batch_size: 16
9   audio:
10    feat_type: 'fbank'
11    feat_dim: 40
12    frame_length: 25
13    frame_shift: 10
14    dither: 0
15    apply_cmvn: True
16    delta_order: 2
17    delta_window_size: 2
18  text:
19    mode: 'character'
20    vocab_file: 'bopomo_vocab_file'
```

TODO: Write your own config file (2/3)

- Training Hyperparams
 - set valid_step to 500 (step for one epoch)
 - max_step: 12k step is enough actually
 - teather forcing: always use it
=> tf_start: 1.0, tf_end:1.0
 - optimizer and lr and eps:
just follow the original settings

```
22  hparas: .....
23      · valid_step: 500
24      · max_step: 12001
25      · tf_start: 1.0
26      · tf_end: 1.0
27      · tf_step: 500000
28      · optimizer: 'Adadelta'
29      · lr: 1.0
30      · eps: 0.00000001 .....
31      · lr_scheduler: 'fixed'
32      · curriculum: 0
```


TODO: Write your own config file (3/3)

- Model Architecture:
 - A thinner model (1~2 layer LSTM) is enough
 - ctc_weight:
 - set to 0.0 to train seq2seq without CTC
 - set between 0.0~1.0 to jointly optimize for CTC + seq2seq
 - Other settings: just follow the original settings

```
34 model: .....
35   ··· ctc_weight: 0.0 .....
36   ··· encoder:
37     ··· prenet: 'vgg' .....
38     ··· # vgg: True .....
39     ··· module: 'LSTM' .....
40     ··· bidirection: True .....
41     ··· dim: [512,512] .....
42     ··· dropout: [0,0] .....
43     ··· layer_norm: [False,False] .....
44     ··· proj: [True,True] .....
45     ··· sample_rate: [1,1] .....
46     ··· sample_style: 'drop' .....
47   ··· attention:
48     ··· mode: 'loc' .....
49     ··· dim: 300 .....
50     ··· num_head: 1 .....
51     ··· v_proj: False .....
52     ··· temperature: 0.5 .....
53     ··· loc_kernel_size: 100 .....
54     ··· loc_kernel_num: 10 .....
55   ··· decoder:
56     ··· module: 'LSTM' .....
57     ··· dim: 512 .....
58     ··· layer: 1 .....
59     ··· dropout: 0 .....
```

TODO: Write your own config file for LM

- See example [config/libri/lm_example.yaml](#)
- Need to modify:
 - path
 - train/dev_split
 - vocab_file
- You can also prepare your own
ㄅ ㄆ ㄇ ㄏ corpus larger much larger
than training set!
 - just replace ['train'] with ['xxx.txt']
where xxx.txt is your collected corpus

```
1  data:
2    corpus: .....
3    # The following depends on corpus
4    name: 'dlhlp' .....
5    path: 'data/DLHLP'
6    train_split: ['train'] .....
7    dev_split: ['dev']
8    bucketing: True
9    batch_size: 32
10   text:
11     mode: 'character' .....
12     vocab_file: 'bopomo_vocab_file'
13
```

Train model!

ASR:

```
$ python3 main.py --config config/dlhelp/asr_dlhelp.yaml
```

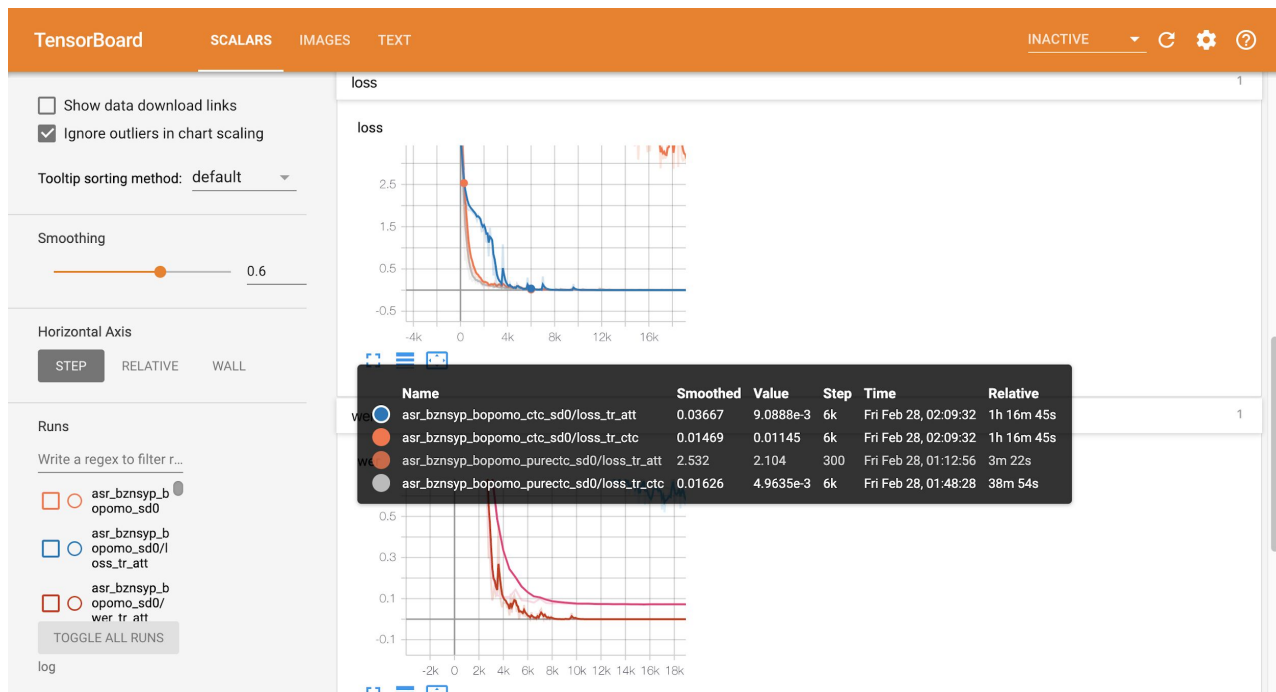
LM:

```
$ python3 main.py --config config/dlhelp/lm_dlhelp.yaml --lm
```

Tensorboard

```
$ tensorboard --logdir log/ --port <port_you_want>
```

Training curve
(loss & WER / train&dev)



Alignment

SCALARS

IMAGES


TEXT

att_align0

att_align0

step **48,500** Fri Feb 28 2020 10:12:30 GMT+0800 (台北標準時間)


asr_bznsyp_bopomo_ctc_sdu



att_align0

step **90,000** Thu Feb 27 2020 10:06:02 GMT+0800 (台北標準時間)

asr_transro_ctc_sdu



SCALARS IMAGES TEXT

2

asr_bznsyp_bopomo_ctc_sd0

丁一么`《X 3-九`尸X 丑 去Y 4 丑` 勿丑` X` 5-4 尸X 尸X` 去丑` 5-世` 尸么` 去丑` 5-世` 尸么`

丁一么'《X 3-九'尸X 乙 去Y 4 口 乙' 勿 乙' X' 乙- 乙 尸X 尸X' 去 乙' 乙- 世' 尸么' 去 乙' 乙- 世' 尸么'

丁一幺`《X 3-九`尸X ㄟ 去Y 4Uㄟ`ㄋㄟ`X`ㄣ-ㄣ 尸X 尸X`去ㄟ`ㄣ-ㄟ`尸幺`去ㄟ`ㄣ-ㄟ`尸幺`

asr_bznsyp_bopomo_ctc_sd0

卅九尤一尤、丁一尤一厂Xㄟ、ㄅY Y、ㄗㄗㄗ、尸、ㄅㄛ、尸、尸尸尸、尸尤、ㄅ一又、丁丁一ㄅ、尸
 XXㄗ 尸尸、〇〇YY

卅九尤一尤、丁一尤、尸Xㄟ、勿Y Y'、卅卅L L'、尸''、么么' P'、P P P'、尸尤、卍一又、T T一么、尸
 X X X' P P P'、口口Y Y'

Test your model - seq2seq

- See example [config/libri/decode_example.yaml](#)
- set max_len_ratio: 0.30
- set ctc_weight: 0.0
- set beam size 2~20 (1 not support)

If you want to use external LM:

- set lm_config&lm_path
- set lm_weight>0.0

```
1  # Most of the parameters will be imported from t
2  src:
3    ckpt: 'ckpt/asr_dlhlp_bopomo_sd0/best_att.pth'
4    config: 'config/dlhlp/asr_dlnlp_bopomo.yaml'
5  data:
6    corpus:
7      name: 'Dlhlp'
8      dev_split: ['dev']
9      test_split: ['test']
10 decode:
11   beam_size: 2
12   min_len_ratio: 0.01
13   max_len_ratio: 0.30
14   lm_path: 'lm_dlhlp_sd0/best_ppx.pth'
15   lm_config: 'config/dlhlp/lm_dlhlp.yaml'
16   lm_weight: 0.3
17   ctc_weight: 0.0
```

Test your model - CTC

- Original code has some problems for jointly decode by seq2seq&CTC
- Add few lines of code to support decode only by CTC (works for ctc_weight=1.0)
- [src/decode.py](#) line 101

```
95     .... # CTC decoding
96     .... if self.apply_ctc:
97     ....     ctc_output = F.log_softmax(
98     ....         self.asr.ctc_layer(encode_feature), dim=-1)
99     ....     ctc_prefix = CTCPrefixScore(ctc_output)
100    ....     ctc_state = ctc_prefix.init_state()
101    ....     if self.ctc_w == 1.0:
102    ....         output_seq = ctc_output[0].argmax(dim=-1)
103    ....         hypothesis = [Hypothesis(decoder_state=dec_state, output_seq=output_seq,
104    ....             output_scores=[0]*len(output_seq), lm_state=None, ctc_prob=0,
105    ....             ctc_state=ctc_state, att_map=None)]
106    ....     return hypothesis
```

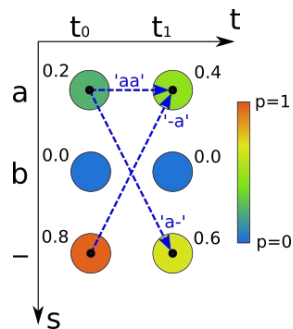

Patch for pure CTC decode (add to src/decode.py line 101)

Simply Greedy Decode!

```
if self.ctc_w == 1.0:
    output_seq = ctc_output[0].argmax(dim=-1)
    hypothesis = [Hypothesis(decoder_state=dec_state, output_seq=output_seq,
                             output_scores=[0]*len(output_seq), lm_state=None, ctc_prob=0,
                             ctc_state=ctc_state, att_map=None)]
    return hypothesis
```

Actually, pure CTC decode also needs beam search.
It could be a good bonus topic for you to implement it.

<https://towardsdatascience.com/beam-search-decoding-in-ctc-trained-neural-networks-5a889a3d85a7>



Test your model - CTC

- set `ctc_weight = 1.0`
- external LM is not supported to jointly decode with CTC in our code
(*But you can implement it as bonus!*)

```
1  # Most of the parameters will be imported from the
2  src:
3  |  · ckpt: 'ckpt/asr_dlhlp_bopomo_sd0/best_ctc.pth'
4  |  · config: 'config/dlhlp/asr_dlnlp_bopomo.yaml'
5  data:
6  |  · corpus:
7  |  |  · name: 'Dlhlp'
8  |  |  · dev_split: ['dev']
9  |  |  · test_split: ['test']
10 decode:
11 |  · beam_size: 2
12 |  · min_len_ratio: 0.01
13 |  · max_len_ratio: 0.30
14 |  · lm_path:
15 |  · lm_config:
16 |  · lm_weight: 0.0
17 |  · ctc_weight: 1.0
```

Test!

```
$ python3 main.py --config <config file> --test --njobs 8
```

- result will be produced at:
 - result/<config file>_dev_output.csv
 - result/<config file>_dev_beam.csv 用不到
 - result/<config file>_test_output.csv
 - result/<config file>_test_beam.csv 用不到
- The format of output file is ``<id> <predicted seq> <truth seq>`` (line by line)
- The truth sequence of testing set was replaced with ㄅ ㄆ ㄇ ㄏ.
- `--njobs` decides the number of threads used for decoding, very important in terms of efficiency. You can set it higher as your machine have more cores.

Process the output from CTC

CTC has repeated tokens

Process it by yourself!

[illegible]

Or you can add additional code to the patch in page 25 to process the 'output_seq' directly before it is written to file.

Eval your output (dev set)

```
$ python3 eval.py --file result/<config_file>_dev_output.csv
```

hint: python2 is not supported here (unicode would be treat as two char)

```
===== Result of result/decode_bopomofo_beam2_lm0.3_dev_output.csv =====
```

Statics	Truth	Prediction	Abs. Diff.
Avg. # of chars	66.99	66.96	0.37
Avg. # of words	17.14	17.12	0.02

report this!

Error Rate (%)	Mean	Std.	Min./Max.
Character	2.1157	2.43	0.00/33.33
Word	7.0030	7.48	0.00/100.00

Note : If the text unit is phoneme, WER = PER and CER is meaningless.

Submit your result (test set)

- Extract the result in kaggle format
- Sample script:
<https://github.com/DLHLP2020/hw1-speech-recognition/blob/master/format.py>
- usage:
`python3 format.py result/<config>_test_output.csv answer.csv`
- upload answer.csv to kaggle

Kaggle rules

- website: <https://www.kaggle.com/c/dlhlp2020spring-asr>
- Your team name should be in [team_github_id]_[any_string]
e.g. daikin_大金
- 5 submission per team & per day
- Using any extra kaggle account to submit is cheating!

Kaggle Evaluation Metric

Mean **Levenshtein Distance** calculated for each sentence in char-level.

- the minimum number of single-character edits (insertions, deletions or substitutions) required to change one word into the other.

$$LevDistance = N_{ins} + N_{del} + N_{sub}$$

Word Error Rate:

$$WER = \frac{N_{ins} + N_{del} + N_{sub}}{N_{target-length}}$$

Baselines (5 points)

Public Simple Baseline: 1.394 (2 pt)

Public Strong Baseline: 1.278 (1 pt)

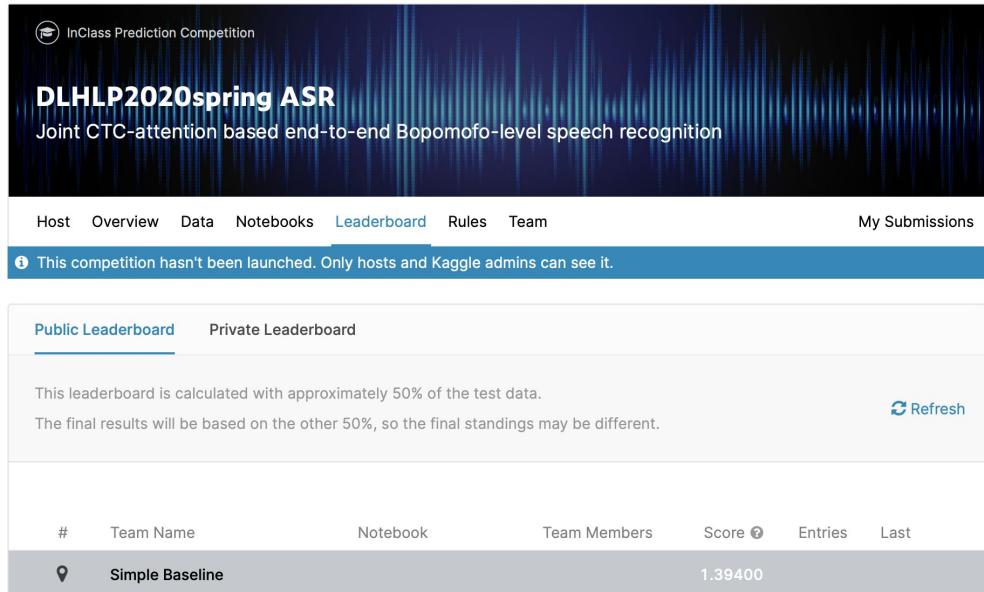
Private Simple Baseline: (?) (1 pt)

Private Strong Baseline: (?) (1 pt)

p.s. private score will be shown
after kaggle deadline

Both would be easy to beat~

<https://www.kaggle.com/c/dlhlp2020spring-asr>



InClass Prediction Competition

DLHLP2020spring ASR

Joint CTC-attention based end-to-end Bopomofo-level speech recognition

Host Overview Data Notebooks **Leaderboard** Rules Team My Submissions

This competition hasn't been launched. Only hosts and Kaggle admins can see it.

Public Leaderboard Private Leaderboard

This leaderboard is calculated with approximately 50% of the test data.
The final results will be based on the other 50%, so the final standings may be different. [Refresh](#)

#	Team Name	Notebook	Team Members	Score ?	Entries	Last
📍	Simple Baseline			1.39400		

Report Questions (1/2)

In 1. ~ 3., just decode without language model.

In 1. ~ 4., set beam size = 2 for speeding.

1. Train a **seq2seq attention-based** ASR model. Paste the learning curve and alignment plot from tensorboard. Report the CER/WER of dev set and kaggle score of testing set. (2 points)
2. Repeat 1. by training a **joint CTC-attention** ASR model (decoding with seq2seq decoder). Which model converges faster? Explain why. (2 points)
3. Use the model in 2. to **decode only in CTC** (ctc_weight=1.0). Report the CER/WER of dev set and kaggle score of testing set. Which model performs better in 1. 2. 3.? Explain why. (2 points)

Report Questions (2/2)

4. Train an **external language model**. Use it to help the model in 1. to decode. Report the CER/WER of dev set and kaggle score of testing set. (2 points)
5. Try decoding the model in 4. with **different beam size** (e.g. 2, 5, 10, 20, 50). Which beam size is the best? (2 points)

Bonus. Other Improvement and Innovation

- If you don't have anything to share, you can just say 'Nothing'. It will be fine.
- Guidance
 - Read new papers
 - Browse github
 - Apply some cool tips

Scoring: Submission

- Kaggle Deadline: 3/22(Sun) 23:59
- Github submission Deadline: 3/25(Wed) before class
- Create a folder 'hw1' under your team Github repo
- 'hw1/' contains:
 - report.pdf
 - reproduce.sh
 - other files and directories
- Report Template: <https://docs.google.com/document/d/1NylgXrlai9Lgysplh9p742zgn2j3cE3B7epqiLDAesE/>
- We restrict Python version 3.6.8 and must compatible with [these package](#)
- Scoring
 - report 10pts
 - kaggle 5pts (over baseline + successfully reproduce)

Github maximum capacity

- within 100MB
- use Dropbox to put your model, use 'wget' to download
- Dropbox Tutorial:
https://docs.google.com/presentation/d/1Sslelij9ZOEN_TGdbAS1oWcl6bT1uSTI6b5_u2wdDc/edit?usp=sharing
- your shell script files should be able to download the model automatically

Scoring: Reproduce

- `bash reproduce.sh $1 $2`
 - ``$1`` is the audio dataset directory (e.g. `data/DLHLP`)
 - ``$2`` is name of the output prediction csv file (e.g. `ans.csv`)
- This script should produce the same result as your **best submission on public leaderboard**.
- Hint: use `sed` to modify the data path (e.g. `'data/DLHLP'`) in your config
- It HAVE to automatically download EVERYTHING you want to wget
- If your code reproduce fail, your **CANNOT** get kaggle score (5%)
- Prepare it carefully!

Scoring: Bonus

- Presentation in class
- Selection criteria
 1. Innovation
 2. Different ways compare to LAS
 - Exiting github is valid, but you have to understand and explain it.
 3. Good performance
- 1 extra pts
- The team quota and the presentation time will be announced based on the time we have.

Late submission policy

- You can submit file until 3 days after deadline
- The score will be calculated as:

$$\text{score}_{\text{final}}(\text{hr}) = \begin{cases} \text{score}_{\text{original}} \times 0.985^{\text{hr}} & , \text{hr} \leq 72 \\ 0 & , \text{hr} > 72 \end{cases}$$

- Late submission form would be announced after deadline



FB Group:
Deep Learning for Human
Language Processing
(2020, Spring)

Q&A

dlhlp.ta@gmail.com